

ECE 571 – Advanced Microprocessor-Based Design Lecture 9

Vince Weaver

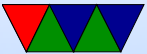
<http://www.eece.maine.edu/~vweaver>

vincent.weaver@maine.edu

16 February 2016

Announcements

- HW#4 Posted, Branch prediction



Notes about HW#3

- Energy

	sleep	stream	MMM	iozone
cores	0.05J	125.8J	20.17J	3.94J
gpu	0J	0J	0J	0J
package	30.8J	211J	30J	31.5J
dram	6.52J	25.4J	1.17J	6.18J
time	10s	9.5s	1.3s	8.38s

- Power



	sleep	stream	MMM	iozone
cores	0.005W	13.2W	15.5W	0.47W
gpu	0W	0J	0J	0J
package	3.1W	22.2W	23.0W	3.76W
dram	0.65W	2.67W	0.9W	0.74W
time	10s	9.5s	1.3s	8.38s

What does the various things exercise?
sleep? stream? MMM? iozone?

- It's a i7-4770, 84W TDP, 22nm, 4-core (8 thread)
- equake_l



threads	time	E	ED	ED2
1	168.5	2836	476k	80M
2	137.1	3155	432k	59M
4	131.8	4174	550k	72M
8	134.2	4538	608k	81M

fastest run time? Some have it with 4 vs 8. Think it's 4 cores, 8 threads.

interesting with values so close, some people got 2 instead of 1 on best ED.

- Why doesn't it scale better? Get whole thesis out of it.

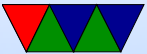


Would think it would.

How would you find out why it doesn't scale?



Oh No, More Caches!



Other Cache Types

- Victim Cache – store last few evicted blocks in case brought back in again, mitigate smaller associativity
- Assist Cache – prefetch into small cache, avoid problem where prefetch kicks out good values
- Trace Cache – store predecoded program traces instead of (or in addition to) instruction cache



Virtual vs Physical Addressing

Programs operate on Virtual addresses.

- PIPT, PIVT (Physical Index, Physical/Virt Tagged) – easiest but requires TLB lookup to translate in critical path
- VIPT, VIVT (Virtual Index, Physical/Virt Tagged) – No need for TLB lookup, but can have aliasing between processes. Can use page coloring, OS support, or ASID (address space id) to keep things separate



Cache Miss Types

- Compulsory (Cold) — miss because first time seen
- Capacity — wouldn't have been a miss with larger cache
- Conflict — miss caused by conflict with another address (would not have been miss with fully assoc cache)
- Coherence — miss caused by other processor



Fixing Compulsory Misses

Prefetching

- Hardware Prefetchers – very good on modern machines. Automatically bring in nearby cachelines.
- Software – loading values before needed also special instructions available
- Large-blocksize of caches. A load brings in all nearby values in the rest of the block.



Fixing Capacity Misses

- Build Bigger Caches



Fixing Conflict Misses

- More Ways in Cache
- Victim Cache
- Code/Variable Alignment, Cache Conscious Data Placement



Fixing Coherence Misses

- False Sharing – independent values in a cache line being accessed by multiple cores



Cache Example 1 – Finding the Parameters

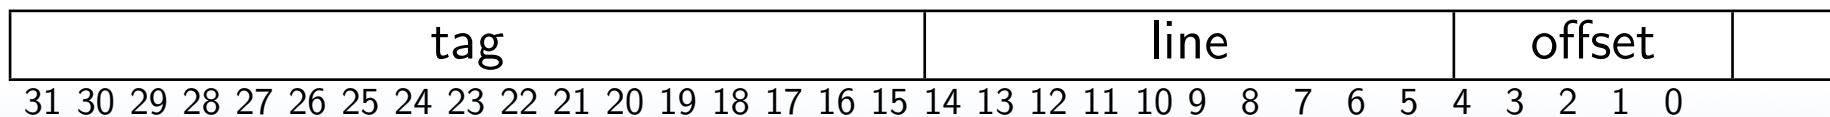
32kB cache (2^{15}), direct mapped (2^0)

32 Byte linesize (2^5), 32-bit address size (2^{32})

offset = $\log_2(\text{linesize}) = 5$ bits

lines = $\log_2((\text{cachesize}/\text{ways})/\text{linesize}) = 1024$ lines
(10 bits)

tag = addresssize - (offset bits + line bits) = 17 bits



Cache Example 2 – Finding the Parameters

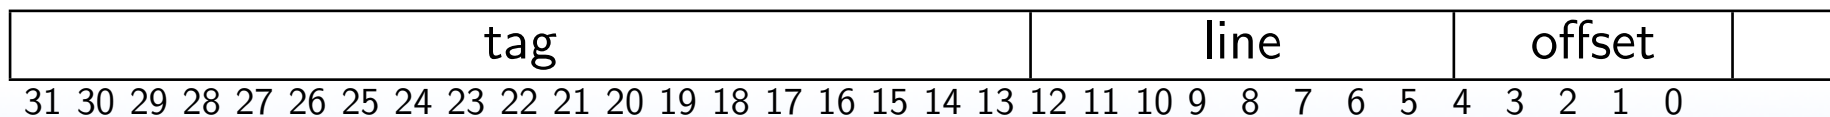
32kB cache (2^{15}), 4-way (2^2)

32 Byte linesize (2^5), 32-bit address size (2^{32})

offset = $\log_2(\text{linesize}) = 5$ bits

lines = $\log_2((\text{cachesize}/\text{ways})/\text{linesize}) = 256$ lines (8 bits)

tag = addresssize - (offset bits + line bits) = 19 bits



Cache Example

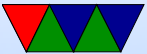
512 Byte cache, 2-Way Set Associative, with 16 byte lines, LRU replacement.

24-bit tag, 16 lines (4 bits), 4-bit offset.

tag																line				offset												
31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	



Cache Example 1



Cache Example – Instruction 1

ldb r1, 0x00000000

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	0	0000 00	0				
1	0				0				
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold



Cache Example – Instruction 2

ldb r1, 0x00000001

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	0	0000 00	0				
1	0				0				
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Hit



Cache Example – Instruction 3

ldb r1, 0x00000010

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	0	0000 00	0				
1	1	0	0	0000 00	0				
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold



Cache Example – Instruction 4

ldb r1, 0x80000010

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	0	0000 00	0				
1	1	0	1	0000 00	1	0	0	8000 00	
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold



Cache Example – Instruction 5

ldb r1, 0xC0000010

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	0	0000 00	0				
1	1	0	0	C000 00	1	0	1	8000 00	
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold (never in cache previously)



Cache Example – Instruction 6

ldb r1, 0xC0000002

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	1	0000 00	1	0	0	c000 00	
1	1	0	0	C000 00	1	0	1	8000 00	
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold



Cache Example – Instruction 7

ldb r1, 0x00000010

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	1	0000 00	1	0	0	c000 00	
1	1	0	1	C000 00	1	0	0	0000 00	
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Conflict



Cache Example – Instruction 8

stb r1, 0x00000005

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	1	0	0000 00	1	0	1	c000 00	
1	1	0	1	C000 00	1	0	0	0000 00	
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Hit



Cache Example 2



Cache Example – Instruction 1

ldrb r1, 0x00000000

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	0	0000 00	0				
1	0				0				
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold



Cache Example – Instruction 2

`str r1, 0x00000001`

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	0	0000 00	0				
1	0				0				
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Hit



Cache Example – Instruction 3

strb r1, 0x00000105

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	1	0000 00	1	1	0	0000 01	
1	0				0				
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold

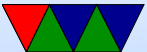


Cache Example – Instruction 4

ldr r1, 0x00000206

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	0	0000 02	1	1	1	0000 01	
1	0				0				
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold



Cache Example – Instruction 5

ldb r1, 0x00000000

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	1	0000 02	1	0	0	0000 00	
1	0				0				
2	0				0				
3	0				0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Conflict



Cache Example – Instruction 6

ldb r1, 0x00000030

Way 0					Way 1				
line	V	D	LRU	Tag	V	D	LRU	Tag	
0	1	0	1	0000 02	1	0	0	0000 00	
1	0				0				
2	0				0				
3	1	0	0	0000 00	0				
4	0				0				
5	0				0				
...									
b	0				0				
c	0				0				
d	0				0				
e	0				0				
f	0				0				

Miss, Cold



CMP Issues

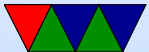
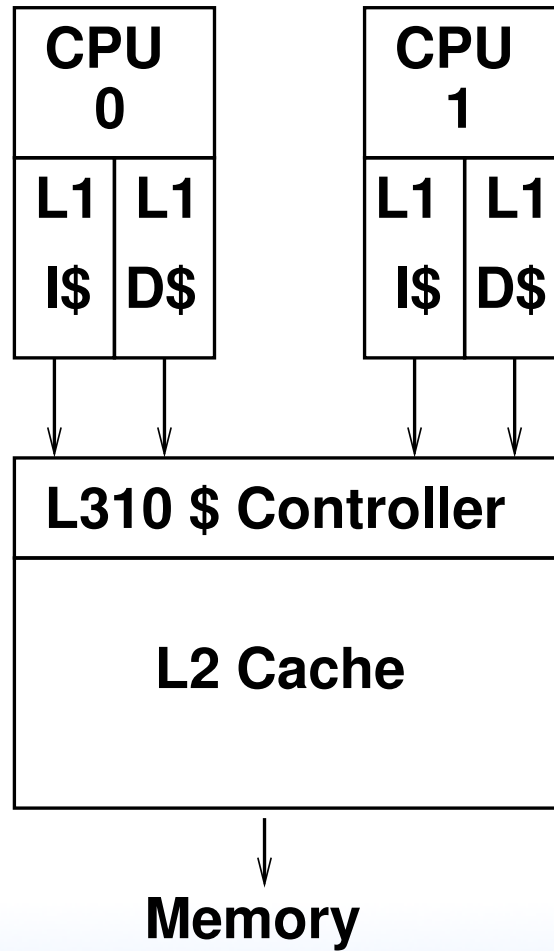


Cache Coherency

- Protocols such as MESI (Modified, Exclusive, Shared, Invalid)
- Snoopy vs Directory



Cortex A9 Cache Layout



Cortex A9 Cache Layout

- OMAP4430 processor
- 32kB 4-way associative, separate L1-I and L1-D
pseudo-round-robin or pseudo random replacement
8-word line size (32B)
critical-word first filling
instruction: VIPT, data: PIPT
- Optional L2 cache controller
Pandaboard has L310 L2 cache controller, 1MB 16-way

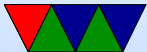
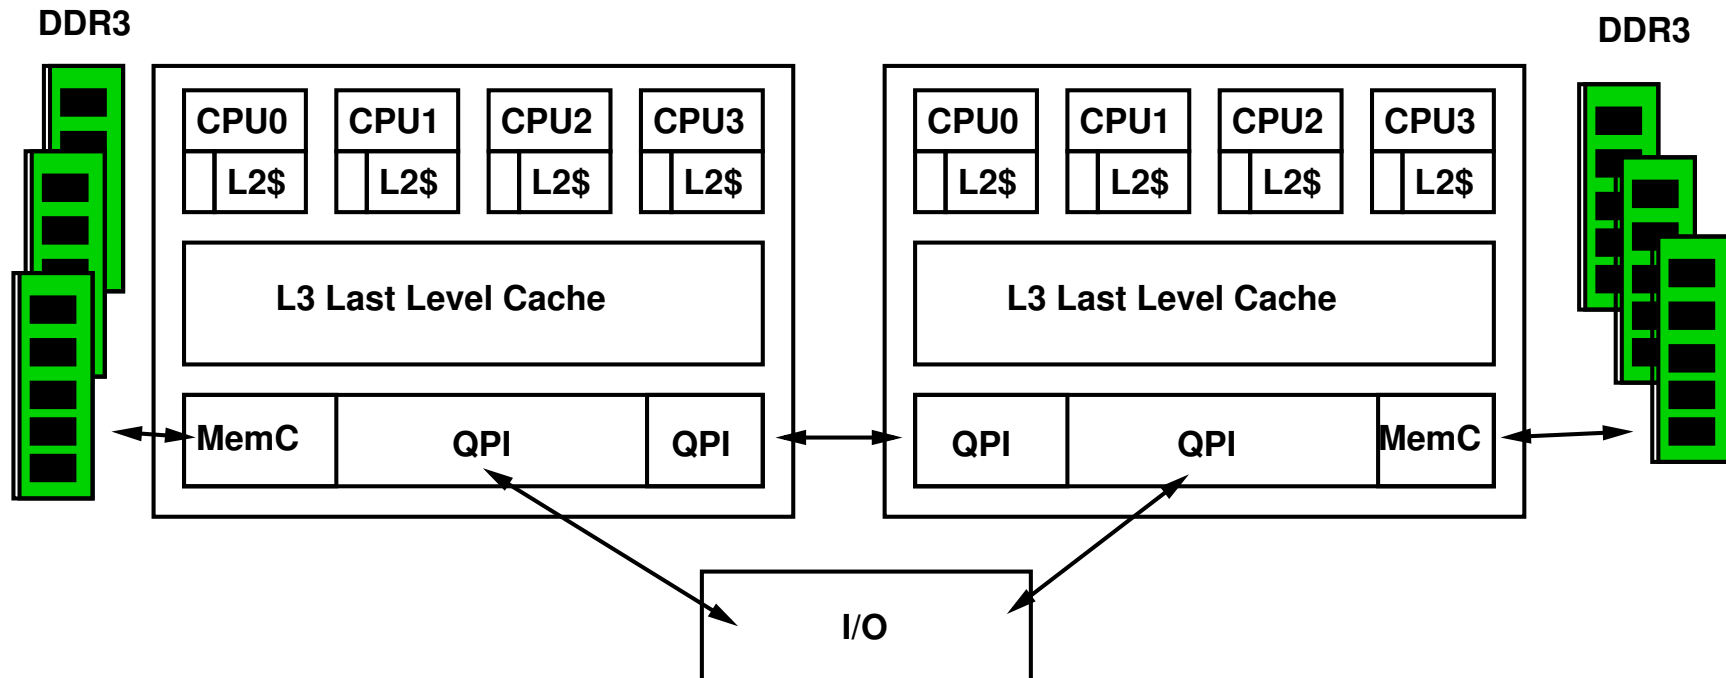


Optional prefetcher

- data cache reads/writes non-blocking, 4 outstanding misses
write buffer: 4 64-bit, allowing write combining



SandyBridge Cache Layout



SandyBridge Cache Layout

- per core 32kB L1 I/D – 4 clocks
64B/line, 8-Way
(shared if hyper-threaded)
writeback
- mOp cache? 1.5K instructions, 8-way, 6Mop/line
Loop stream detector, can execute w/o accessing icache
- per core 256kB L2 unified – 12 clocks



64B/line, 8-way
writeback. non-inclusive

- shared L3 1MB-20MB – 26-31 clocks
64B/line. 12-way (varies)
writeback, inclusive
- various hw prefetchers operating



Cache Performance Measurement

Matrix-Matrix multiply is the typical example.

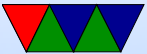
Despite being a big deal in HPC, MMM happens in embedded world too.



Naive Matrix-Matrix Multiply 1

```
#define MATRIX_SIZE 512
static double a[MATRIX_SIZE][MATRIX_SIZE];
static double b[MATRIX_SIZE][MATRIX_SIZE];
static double c[MATRIX_SIZE][MATRIX_SIZE];

for(j=0; j<MATRIX_SIZE; j++) {
    for(i=0; i<MATRIX_SIZE; i++) {
        for(k=0; k<MATRIX_SIZE; k++) {
            c[i][j] += a[i][k] * b[k][j];
        }
    }
}
```



Naive Matrix-Matrix Multiply 1 – what's the issue?

- Branch Misses?
- TLB Misses?
- ICache Misses?
- DCache Misses?
- L2 Cache Misses?



Naive Matrix-Matrix Multiply 1 – perf results



Matrix multiply sum: s=27665734022509.746094

Performance counter stats for './matrix_multiply':

11296.203614	task-clock	#	0.999 CPUs utilized
20	context-switches	#	0.000 M/sec
0	CPU-migrations	#	0.000 M/sec
1,633	page-faults	#	0.000 M/sec
9,032,356,979	cycles	#	0.800 GHz
6,547,102	stalled-cycles-frontend	#	0.07% frontend cycles id
8,213,005,758	stalled-cycles-backend	#	90.93% backend cycles id
1,176,144,886	instructions	#	0.13 insns per cycle
		#	6.98 stalled cycles per
137,651,296	branches	#	12.186 M/sec
795,064	branch-misses	#	0.58% of all branches
11.303802490	seconds time elapsed		



Naive Matrix-Matrix Multiply 1 – DCache results

```
vince@arm:~/class/ece571/lecture10_code$ perf stat -e cache-misses,cache-r  
Matrix multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply':
```

```
171,786,441 cache-misses          # 42.072 % of all cache ref  
408,318,876 cache-references
```

```
10.680664062 seconds time elapsed
```



Naive Matrix-Matrix Multiply 1 – ICache results

```
vince@arm:~/class/ece571/lecture10_code$ perf stat -e l1-icache-load-misses  
Matrix multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply':
```

```
          536,719 l1-icache-load-misses  
1,174,927,869 instructions                #    0.00  insns per cycle
```

```
12.203002930 seconds time elapsed
```

```
0.04% icache misses
```



Naive Matrix-Matrix Multiply 1 – TLB Misses

```
vince@arm:~/class/ece571/lecture10_code$ perf stat -e dTLB-load-misses,dTLB-store-misses ./matrix_multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply':
```

```
135,253,464 dTLB-load-misses
```

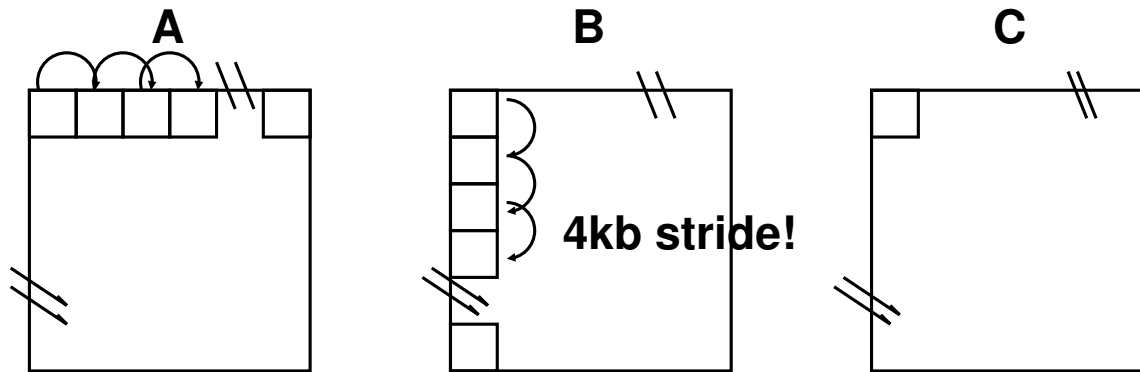
```
135,253,464 dTLB-store-misses
```

```
12.443572998 seconds time elapsed
```



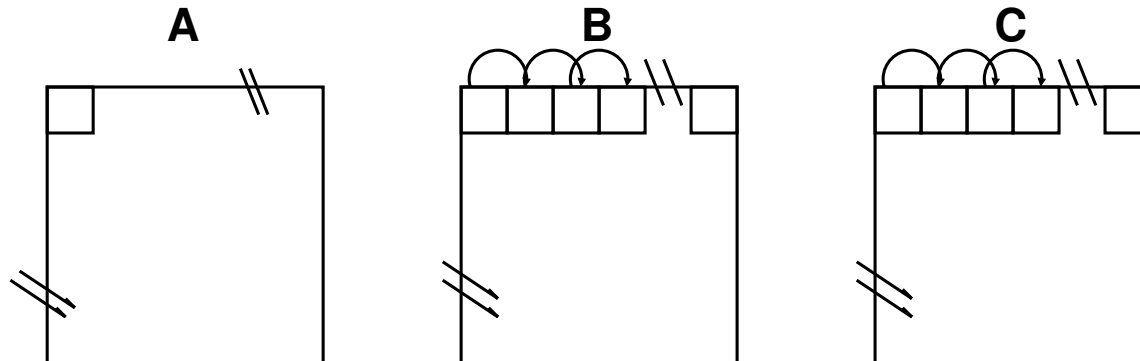
Naive Matrix-Matrix Multiply 1 – What's the Issue?

400M memory accesses about right ($512 \times 512 \times 512 \times 3$)



Switch the loop ordering

```
for (i=0; i<MATRIX_SIZE; i++) {  
    for (k=0; k<MATRIX_SIZE; k++) {  
        for (j=0; j<MATRIX_SIZE; j++) {  
            c[i][j] += a[i][k] * b[k][j];  
        }  
    }  
}
```



Naive Matrix-Matrix Multiply 2 – perf results



Matrix multiply sum: s=27665734022509.746094

Performance counter stats for './matrix_multiply_swapped':

3443.267822	task-clock	#	0.999 CPUs utilized
5	context-switches	#	0.000 M/sec
0	CPU-migrations	#	0.000 M/sec
1,633	page-faults	#	0.000 M/sec
2,849,573,010	cycles	#	0.828 GHz
2,913,607	stalled-cycles-frontend	#	0.10% frontend cycles id
1,893,138,507	stalled-cycles-backend	#	66.44% backend cycles id
965,962,767	instructions	#	0.34 insns per cycle
		#	1.96 stalled cycles per
136,649,964	branches	#	39.686 M/sec
553,643	branch-misses	#	0.41% of all branches
3.447875977	seconds time elapsed		



Naive Matrix-Matrix Multiply 2 – DCache results

```
vince@arm:~/class/ece571/lecture10_code$ perf stat -e cache-misses,cache-r  
Matrix multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply_swapped':
```

```
      38,628,043 cache-misses          #    9.409 % of all cache ref  
    410,528,663 cache-references  
  
    5.585754395 seconds time elapsed
```



Naive Matrix-Matrix Multiply 2 – ICache results

```
vince@arm:~/class/ece571/lecture10_code$ perf stat -e ll-icache-load-misses  
Matrix multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply_swapped':
```

```
      254,041 ll-icache-load-misses  
964,335,795 instructions          #    0.00  insns per cycle
```

```
4.245208740 seconds time elapsed
```

```
0.02%
```



Naive Matrix-Matrix Multiply 1 – TLB Misses

```
vince@arm:~/class/ece571/lecture10_code$ perf stat -e dTLB-load-misses,dTLB-store-misses ./matrix_multiply_swapped  
Matrix multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply_swapped':
```

```
486,039 dTLB-load-misses
```

```
486,039 dTLB-store-misses
```

```
5.242126465 seconds time elapsed
```



Other Ways to Optimize

- Tiling
- Parallelizing



Use ATLAS/BLAS

```
cbblas_dgemm(CblasRowMajor ,  
             CblasNoTrans ,CblasNoTrans ,  
             512,512,512 ,  
             1.0,(const double *)a,512 ,  
             (const double *)b,512 ,  
             1.0,(double *)c,512);
```



Matrix-Matrix Mul ATLAS – perf results

```
Matrix multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply_atlas':
```

1158.325193	task-clock	#	1.678 CPUs utilized
12	context-switches	#	0.000 M/sec
1	CPU-migrations	#	0.000 M/sec
2,017	page-faults	#	0.002 M/sec
597,931,712	cycles	#	0.516 GHz
2,043,500	stalled-cycles-frontend	#	0.34% frontend cycles id
258,860,537	stalled-cycles-backend	#	43.29% backend cycles id
519,715,833	instructions	#	0.87 insns per cycle
		#	0.50 stalled cycles per
36,716,368	branches	#	31.698 M/sec
815,440	branch-misses	#	2.22% of all branches

```
0.690429687 seconds time elapsed
```



Matrix-Matrix Mul ATLAS – DCache

```
Matrix multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply_atlas':
```

```
    11,988,047 cache-misses          #    8.128 % of all cache ref  
   147,494,664 cache-references  
  
    0.598632813 seconds time elapsed
```



Matrix-Matrix Mul ATLAS – TLB

```
vince@arm:~/class/ece571/lecture10_code$ perf stat -e dTLB-load-misses ./m  
Matrix multiply sum: s=27665734022509.746094
```

```
Performance counter stats for './matrix_multiply_atlas':
```

```
224,299 dTLB-load-misses
```

```
0.755981446 seconds time elapsed
```

